

# DQN Trading Algorithm: Step-by-Step Execution Flow

## Phase 0: Initialization (Happens Once)

### Step 0.1: Environment Creation

```
python
```

```
cfg = EnvConfig(batch_size=256, T=10)
env = Environment(cfg)
```

#### What happens:

1. Creates 256 parallel trading environments (like 256 traders in parallel universes)
2. Each environment tracks:
  - `self.price`: [256, 10] - last 10 prices for each environment
  - `self.pos`: [256, 10] - last 10 positions for each environment
3. Generates initial price history by simulating mean-reverting process:
  - Start at  $S_m = 100$
  - For 10 time steps: 
$$S(t+1) = S(t) + \sigma * \varepsilon - \kappa * (S(t) - S_m)$$
  - Clamp between 80-120
4. Creates initial state `self.x` [256, 20] by concatenating normalized prices + positions

### Step 0.2: Agent Creation

```
python
```

```
agent = Agent(gamma=0.99, epsilon=1.0, lr=3e-4, input_dims=20, ...)
```

#### What happens:

1. Creates neural network (`Q_eval`):
  - Input: 20 dimensions (10 prices + 10 positions)
  - Hidden: 256 → 256 neurons with ReLU
  - Output: 3 Q-values (one for each action: sell, hold, buy)
2. Initializes Adam optimizer
3. Allocates replay buffer memory:
  - `state_memory`: [100000, 20] - stores past states

- `action_memory`: [100000] - stores past actions
- `reward_memory`: [100000] - stores past rewards
- `new_state_memory`: [100000, 20] - stores next states
- `terminal_memory`: [100000] - stores done flags

4. Sets `mem_cntr = 0` (no memories yet)

5. Sets `epsilon = 1.0` (100% exploration at start)

---

## Phase 1: Training Loop (Repeats for 500 Episodes)

### EPISODE 1, STEP 1

#### Step 1.1: Reset Environment

```
python
state = env.reset() # Returns [256, 20]
```

- Regenerates fresh price histories for all 256 environments
- All positions set to 0
- Returns normalized state tensor

#### Step 1.2: Choose Actions (Epsilon-Greedy)

```
python
actions = agent.choose_action(state) # Returns [256] numpy array
```

#### Detailed process:

1. Convert state to torch tensor and move to device

#### 2. Forward pass through neural network:

Input [256, 20] → FC1 [256, 256] → ReLU → FC2 [256, 256] → ReLU → FC3 [256, 3]  
 Output: Q-values [256, 3]

#### 3. Per-row epsilon-greedy:

- For each of 256 environments independently:
  - Generate random number  $r \in [0,1]$
  - If  $r < \text{epsilon}$  (currently 1.0):

- Choose random action from {0, 1, 2}
- Else:
  - Choose action = argmax(Q-values) for that row
- Initially epsilon=1.0, so ALL 256 environments choose random actions

4. Return numpy array [256] with action indices

**Example:** `[actions = [1, 0, 2, 1, 0, ...]]` (256 actions total)

### Step 1.3: Environment Step

```
python
next_state, rewards, dones, _ = env.step(actions)
```

### Detailed process for EACH of 256 environments:

#### 1. Map action index to actual position:

```
python
a_t = self._idx2val[actions] # [256] in {-1, 0, +1}
```

- action 0 → position -1 (sell)
- action 1 → position 0 (hold)
- action 2 → position +1 (buy)

#### 2. Get current state:

```
python
prev_price = self.price[:, -1]    # [256] - last price
prev_pos = self.pos[:, -1]        # [256] - last position
```

#### 3. Generate next price (mean-reverting process):

```
python
ε ~ N(0,1) # Random shock for each environment
next_price = prev_price + σ*ε - κ*(prev_price - S_m)
next_price = clamp(round(next_price), 80, 120)
```

- Each environment gets different random shock
- Prices independently evolve

#### 4. Calculate reward for each environment:

```
python

dS = next_price - prev_price

pnL = a_t * dS          # P&L from position
friction = β * |a_t - prev_pos| # Cost of changing position
vol_penalty = γ * σ² * a_t²    # Risk penalty

reward = pnL - friction - vol_penalty
```

#### Example for one environment:

- $\text{prev\_price} = 98, \text{next\_price} = 100 \rightarrow dS = +2$
- $a_t = +1$  (buy),  $\text{prev\_pos} = 0$
- $\text{pnL} = (+1) * (+2) = +2.0$
- $\text{friction} = 0.1 * |1 - 0| = 0.1$
- $\text{vol\_penalty} = 0.01 * 4 * 1 = 0.04$
- **reward = 2.0 - 0.1 - 0.04 = 1.86 ✓ Good reward!**

#### 5. Update price and position histories (rolling window):

```
python

self.price = [price[:, 1:], next_price] # Drop oldest, append newest
self.pos = [pos[:, 1:], a_t]           # Drop oldest, append newest
```

- Each is  $[256, 10]$  - sliding window of last 10 values

#### 6. Update state representation:

```
python

z_price = (self.price - 100) / 2.0 # Normalize
self.x = concat([z_price, self.pos], dim=1) # [256, 20]
```

#### 7. Return:

- $\text{next\_state: } [256, 20]$
- rewards:  $[256]$  floats
- dones:  $[256]$  all False (we don't have terminal states)

## Step 1.4: Store Transitions in Replay Buffer

```
python
```

```
agent.store_transition(state, actions, rewards, next_state, dones)
```

### What happens:

1. Convert all torch tensors to numpy arrays
2. Calculate indices for circular buffer:

```
python
```

```
indices = [0, 1, 2, ..., 255] # First 256 slots
```

3. Store in replay buffer:

```
python
```

```
state_memory[0:256] = state      #[256, 20]
action_memory[0:256] = actions    #[256]
reward_memory[0:256] = rewards    #[256]
new_state_memory[0:256] = next_state #[256, 20]
terminal_memory[0:256] = dones     #[256]
```

4. Increment memory counter: `mem_cntr = 256`

## Step 1.5: Learn (First Call - SKIPPED!)

```
python
```

```
agent.learn()
```

### What happens:

```
python
```

```
if self.mem_cntr < self.batch_size: # 256 < 128? No, 256 >= 128
    return # Actually will NOT return, will try to learn
```

Wait, we have 256 memories and batch\_size=128, so we DO learn!

### Learning process (FIRST TIME):

1. Sample random batch from replay buffer:

```
python
```

```
batch = random.choice([0, 1, ..., 255], size=128, replace=False)  
# Example: batch = [45, 203, 12, 189, ...]
```

## 2. Load batch from memory:

```
python
```

```
states = state_memory[batch]    #[128, 20]  
actions = action_memory[batch]  #[128]  
rewards = reward_memory[batch] #[128]  
next_states = new_state_memory[batch] #[128, 20]  
dones = terminal_memory[batch]  #[128]
```

## 3. Forward pass to get Q-values for current states:

```
python
```

```
q_values = Q_eval(states) #[128, 3]  
# Example row: [0.23, -0.15, 0.08] (random initially)
```

## 4. Select Q-values for actions that were actually taken:

```
python
```

```
row_idx = [0, 1, 2, ..., 127]  
q_eval = q_values[row_idx, actions] #[128]  
# If actions[0]=1, picks q_values[0, 1] = -0.15
```

## 5. Compute target Q-values (what we WANT Q to predict):

```
python
```

```
with torch.no_grad(): # Don't backprop through target  
    q_next_all = Q_eval(next_states) #[128, 3]  
    q_next_max = max(q_next_all, dim=1) #[128] - best action  
  
    # Bellman equation:  
    q_target = rewards + γ * (1 - dones) * q_next_max
```

## Example for one sample:

- reward = 1.86
- q\_next\_max = 0.31 (best Q-value for next state)

- $q_{\text{target}} = 1.86 + 0.99 * 0.31 = 2.17$

## 6. Compute loss (Mean Squared Error):

```
python
```

```
loss = MSE(q_eval, q_target)
# Example: MSE([-0.15, 0.23, ...], [2.17, 1.45, ...])
```

## 7. Backpropagation:

```
python
```

```
loss.backward() # Compute gradients ∂loss/∂weights
optimizer.step() # Update weights: w = w - lr * ∇loss
```

### What this does:

- Adjusts all  $20256 + 256256 + 256*3 \approx 70,000$  parameters
- Makes Q\_eval predict values closer to q\_target
- Agent learns: "When I see this state and take this action, I should expect ~2.17 reward"

## 8. Decay epsilon:

```
python
```

```
epsilon = max(0.05, 1.0 - 0.00001)
epsilon = 0.99999 # Barely changed
```

### Step 1.6: Update State

```
python
```

```
state = next_state #[256, 20]
```

## EPISODE 1, STEPS 2-200

### Repeat steps 1.2-1.6 for 199 more times:

- Each step, 256 new transitions are stored → memory fills up
- After 200 steps:  $\boxed{\text{mem\_cntr} = 256 * 200 = 51,200}$  memories
- Agent has learned 200 times
- Epsilon decayed to:  $\boxed{1.0 - 200*0.00001 = 0.998}$

- Q-network weights updated 200 times
  - Agent slowly learns patterns: "If price is high and falling, holding gives better Q-value than buying"
- 

## EPISODES 2-500

Same process repeats:

### Episode 10 (2,000 steps total):

- `mem_cntr = 512,000` but buffer only holds 100,000 → oldest memories overwritten
- $\text{epsilon} = 0.98 \rightarrow 98\% \text{ exploration}, 2\% \text{ exploitation}$
- Q-values starting to make sense:  $Q(\text{good\_state}, \text{good\_action}) > Q(\text{bad\_state}, \text{bad\_action})$

### Episode 100 (20,000 steps):

- $\text{epsilon} = 0.80 \rightarrow 80\% \text{ exploration}, 20\% \text{ exploitation}$
- Agent occasionally picks learned good actions
- Q-network has seen millions of gradient updates

### Episode 250 (50,000 steps):

- $\text{epsilon} = 0.50 \rightarrow \text{balanced exploration/exploitation}$
- Agent frequently picks learned actions
- Rewards starting to improve: was -0.10 avg, now -0.05 avg

### Episode 400 (80,000 steps):

- $\text{epsilon} = 0.20 \rightarrow \text{mostly exploitation}$
- Agent mostly follows learned policy
- Rewards positive: +0.02 avg

### Episode 500 (100,000 steps):

- $\text{epsilon} = 0.05 \rightarrow 95\% \text{ exploitation}, 5\% \text{ exploration}$
  - Agent has converged to learned policy
  - Final avg reward: +0.1265 per step
- 

## Phase 2: Evaluation (Pure Exploitation)

### Step E.1: Disable Exploration

```
python
```

```
agent.epsilon = 0.0 # GREEDY ONLY
```

## Step E.2: Run 20 Evaluation Episodes

For each episode:

### 1. Reset environment

```
python
```

```
state = env.reset()
```

### 2. Run 100 steps with pure exploitation:

```
python
```

```
for step in range(100):
    actions = agent.choose_action(state)
    # With epsilon=0, ALWAYS picks argmax(Q-values)
    # No randomness!

    next_state, rewards, dones, _ = env.step(actions)
    total_reward += rewards.mean()
    state = next_state
```

### 3. Track total reward for episode

- Example: Episode 1 = 16.61, Episode 2 = 16.34, ...

## Step E.3: Compute Statistics

```
python
```

Average: 17.04 ± 0.54

Min: 16.33

Max: 18.06

## What the Agent Actually Learned

By the end, the neural network has learned a **value function**  $Q(s, a)$ :

**Input:** [normalized prices for last 10 steps, positions for last 10 steps]

**Output:**  $[Q(s, \text{sell}), Q(s, \text{hold}), Q(s, \text{buy})]$

## Example learned behaviors:

### 1. Mean reversion:

- If price = 115 (high) and falling →  $Q(s, \text{sell}) = 2.5$  (highest)
- If price = 85 (low) and rising →  $Q(s, \text{buy}) = 2.8$  (highest)

### 2. Friction avoidance:

- If current position = +1 and  $Q(\text{buy})$  only slightly better →  $Q(\text{hold})$  chosen (avoids cost)

### 3. Risk management:

- Large positions penalized → agent prefers smaller positions unless strong signal

The agent essentially learned to be a **statistical arbitrage trader** that:

- Buys when prices are low relative to mean
  - Sells when prices are high relative to mean
  - Avoids excessive trading (friction costs)
  - Manages position risk
- 

## Key Insight: Bootstrapping

The magic of Q-learning is **bootstrapping** - the agent uses its own estimates to improve its own estimates:

1. **Initially:** Q-values are random → agent explores randomly
2. **Early learning:** Agent discovers "buying at 85 gives +2 reward" → updates  $Q(s, \text{buy})$  toward 2
3. **Later:** Agent sees state that leads to state where  $Q(s, \text{buy})=2$  → updates previous state to  $2 + \gamma * 2 = 3.98$
4. **Eventually:** Chain of Q-values propagates backward, agent learns long-term consequences

It's literally "pulling itself up by its bootstraps" - using estimates to improve estimates!